

National University of Technology



Computer Science Department

Semester Spring– 2025

Program: Artificial intelligence

Course: Programming for AI

Course Code: CS282

Assignment - 03

Submitted To:

Umar Aftab

Submitted By:

Muhammad Ahad Imran

F23607034

```
julia> iris = dataset("datasets", "iris")
150x5 DataFrame
  Row   SepalLength   SepalWidth   PetalLength   PetalWidth   Species
      Float64      Float64      Float64      Float64      Cat...
```

Row	SepalLength Float64	SepalWidth Float64	PetalLength Float64	PetalWidth Float64	Species Cat...
1	5.1	3.5	1.4	0.2	setosa
2	4.9	3.0	1.4	0.2	setosa
3	4.7	3.2	1.3	0.2	setosa
4	4.6	3.1	1.5	0.2	setosa
5	5.0	3.6	1.4	0.2	setosa
6	5.4	3.9	1.7	0.4	setosa
7	4.6	3.4	1.4	0.3	setosa
8	5.0	3.4	1.5	0.2	setosa
9	4.4	2.9	1.4	0.2	setosa
10	4.9	3.1	1.5	0.1	setosa
11	5.4	3.7	1.5	0.2	setosa
⋮	⋮	⋮	⋮	⋮	⋮
141	6.7	3.1	5.6	2.4	virginica
142	6.9	3.1	5.1	2.3	virginica
143	5.8	2.7	5.1	1.9	virginica
144	6.8	3.2	5.9	2.3	virginica
145	6.7	3.3	5.7	2.5	virginica
146	6.7	3.0	5.2	2.3	virginica
147	6.3	2.5	5.0	1.9	virginica
148	6.5	3.0	5.2	2.0	virginica
149	6.2	3.4	5.4	2.3	virginica
150	5.9	3.0	5.1	1.8	virginica

129 rows omitted

```
julia> X = Matrix(iris[:, 1:4])' # size: (4, 150)
4x150 adjoint(::Matrix{Float64}) with eltype Float64:
 5.1  4.9  4.7  4.6  5.0  5.4  4.6  5.0  4.4  4.9  5.4  4.8  ...  6.9  6.7  6.9  5.8  6.8  6.7  6.7  6.3  6.5  6.2  5.9
 3.5  3.0  3.2  3.1  3.6  3.9  3.4  3.4  2.9  3.1  3.7  3.4  ...  3.1  3.1  3.1  2.7  3.2  3.3  3.0  2.5  3.0  3.4  3.0
 1.4  1.4  1.3  1.5  1.4  1.7  1.4  1.5  1.4  1.5  1.5  1.6  ...  5.4  5.6  5.1  5.1  5.9  5.7  5.2  5.0  5.2  5.4  5.1
 0.2  0.2  0.2  0.2  0.2  0.4  0.3  0.2  0.2  0.1  0.2  0.2  ...  2.1  2.4  2.3  1.9  2.3  2.5  2.3  1.9  2.0  2.3  1.8
```



```
julia> datasett = [(X[:, i], Y[:, i]) for i in 1:size(X, 2)]
150-element Vector{Tuple{Vector{Float64}, OneHotArrays.OneHotVector{UInt32}}}:
([[-0.8976738791967663, 1.015601990713633, -1.3357516342415194, -1.311052148205131], [1, 0, 0])
([-1.1392004834649534, -0.1315388120502606, -1.3357516342415194, -1.311052148205131], [1, 0, 0])
([-1.3807270877331417, 0.32731750905529733, -1.3923992862449763, -1.311052148205131], [1, 0, 0])
([-1.5014903898672363, 0.09788934850251835, -1.2791039822380623, -1.311052148205131], [1, 0, 0])
([-1.01843718133086, 1.2450301512664121, -1.3357516342415194, -1.311052148205131], [1, 0, 0])
([-0.5353839727944835, 1.9333146329247481, -1.165808678231148, -1.0486667949952981], [1, 0, 0])
([-1.5014903898672363, 0.7861738301608542, -1.3357516342415194, -1.1798594716002144], [1, 0, 0])
([-1.01843718133086, 0.7861738301608542, -1.2791039822380623, -1.311052148205131], [1, 0, 0])
([-1.7430169941354234, -0.36096697260303956, -1.3357516342415194, -1.311052148205131], [1, 0, 0])
([-1.1392004834649534, 0.09788934850251835, -1.2791039822380623, -1.4422448248100472], [1, 0, 0])
([-0.5353839727944835, 1.4744503118191912, -1.2791039822380623, -1.311052148205131], [1, 0, 0])
([-1.2599637855990482, 0.7861738301608542, -1.2224563302346052, -1.311052148205131], [1, 0, 0])
([-1.2599637855990482, -0.1315388120502606, -1.3357516342415194, -1.4422448248100472], [1, 0, 0])
...
([0.18919584001008014, -0.1315388120502606, 0.590268533876023, 0.7880306774735305], [0, 0, 1])
([1.2760655592169265, 0.09788934850251835, 0.9301544458967661, 1.1816087072882795], [0, 0, 1])
([1.0345389549487385, 0.09788934850251835, 1.04344974990368, 1.5751867371030284], [0, 0, 1])
([1.2760655592169265, 0.09788934850251835, 0.7602114898863943, 1.4439940604981119], [0, 0, 1])
([-0.05233076425810807, -0.8198232937085965, 0.7602114898863943, 0.9192233540784467], [0, 0, 1])
```

```
julia> train_data, test_data = splitobs(shuffleobs(datasett), at=0.8)
(Tuple{Vector{Float64}, OneHotArrays.OneHotVector{UInt32}})(([0.18919584001008014, 0.7861738301608542, 0.4203255778656516, 0.5256453242636979], [0, 1, 0]), ([2.2421719762896783, -0.1315388120502606, 1.3266880099209655, 1.4439940604981119], [0, 0, 1]), ([0.18919584001008014, -1.9669640964724904, 0.137087317848366, -0.2615107353658002], [0, 1, 0]), ([1.0345389549487385, 0.5567456696080753, 1.1000974019071375, 1.1816087072882795], [0, 0, 1]), ([0.8976738791967663, 1.015601990713633, -1.3357516342415194, -1.311052148205131], [1, 0, 0]), ([0.7930123506805502, -0.1315388120502606, 0.986802097900223, 0.7880306774735305], [0, 0, 1]), ([0.05233076425810807, -0.8198232937085965, 0.19373496985182292, -0.2615107353658002], [0, 1, 0]), ([0.06843253787598658, 0.32731750905529733, 0.590268533876023, 0.7880306774735305], [0, 1, 0]), ([0.18919584001008014, -0.36096697260303956, 0.42032557786565167, 0.39445264765878146], [0, 1, 0]), ([0.5514857464123619, -0.8198232937085965, 0.6469161858794804, 0.7880306774735305], [0, 0, 1]), ([0.5353839727944835, 1.9333146329247481, -1.3923992862449763, -1.0486667949952981], [1, 0, 0]), ([0.7430169941354234, -0.36096697260303956, -1.3357516342415194, -1.311052148205131], [1, 0, 0]), ([0.01843718133086, -2.4258204175780484, -0.14615094216891966, -0.2615107353658002], [0, 1, 0]), ([0.05233076425810807, -0.8198232937085965, 0.7602114898863943, 0.9192233540784467], [0, 0, 1]), ([2.4836985805578666, 1.703886472371969, 1.4966309659313373, 1.0504160306833632], [0, 0, 1]), ([0.9137756528146437, -0.1315388120502606, 0.3636779258621947, 0.263259971053865], [0, 1, 0]), ([0.18919584001008014, -0.1315388120502606, 0.590268533876023, 0.7880306774735305], [0, 0, 1]), ([1.0345389549487385, -0.1315388120502606, 0.7035638378829373, 0.6568380088686141], [0, 1, 0])
```

```
julia> model = Chain(
    Dense(4, 16, relu),
    Dense(16, 8, relu),
    Dense(8, 3),
    softmax
)

Chain(
  Dense(4 => 16, relu), # 80 parameters
  Dense(16 => 8, relu), # 136 parameters
  Dense(8 => 3), # 27 parameters
  softmax,
) # Total: 6 arrays, 243 parameters, 1.254 KiB.
```

```
julia> loss(x, y) = Flux.crossentropy(model(x), y)
loss (generic function with 1 method)
```

```
julia> accuracy(x, y) = mean(Flux.onecold(model(x)) .== Flux.onecold(y))
accuracy (generic function with 1 method)
```

[illegible]

```
julia> epochs = 50
50
```

```
julia> train_loss = Float64[]
Float64[]
```

```
julia> train_acc = Float64[]
Float64[]
```

```
julia> for epoch in 1:epochs
    for (x, y) in train_data
        x = Float32(x) # Ensure correct type
        y = Float32(y)
        gs = Flux.gradient(model) do m
            loss(x, y)
        end
        Flux.update!(opt_state, model, gs[1])
    end
    l = loss(hcat(first.(train_data)...), hcat(last.(train_data)...))
    a = accuracy(hcat(first.(train_data)...), hcat(last.(train_data)...))
    push!(train_loss, l)
    push!(train_acc, a)
    println("Epoch $epoch - Loss: $(round(l, digits=4)) | Accuracy: $(round(a * 100, digits=2))%")
end
```


```
Epoch 1 - Loss: 1.078 | Accuracy: 38.33%
Epoch 2 - Loss: 1.078 | Accuracy: 38.33%
Epoch 3 - Loss: 1.078 | Accuracy: 38.33%
Epoch 4 - Loss: 1.078 | Accuracy: 38.33%
Epoch 5 - Loss: 1.078 | Accuracy: 38.33%
Epoch 6 - Loss: 1.078 | Accuracy: 38.33%
Epoch 7 - Loss: 1.078 | Accuracy: 38.33%
Epoch 8 - Loss: 1.078 | Accuracy: 38.33%
Epoch 9 - Loss: 1.078 | Accuracy: 38.33%
Epoch 10 - Loss: 1.078 | Accuracy: 38.33%
Epoch 11 - Loss: 1.078 | Accuracy: 38.33%
Epoch 12 - Loss: 1.078 | Accuracy: 38.33%
Epoch 13 - Loss: 1.078 | Accuracy: 38.33%
Epoch 14 - Loss: 1.078 | Accuracy: 38.33%
Epoch 15 - Loss: 1.078 | Accuracy: 38.33%
Epoch 16 - Loss: 1.078 | Accuracy: 38.33%
Epoch 17 - Loss: 1.078 | Accuracy: 38.33%
Epoch 18 - Loss: 1.078 | Accuracy: 38.33%
Epoch 19 - Loss: 1.078 | Accuracy: 38.33%
Epoch 20 - Loss: 1.078 | Accuracy: 38.33%
Epoch 21 - Loss: 1.078 | Accuracy: 38.33%
Epoch 22 - Loss: 1.078 | Accuracy: 38.33%
Epoch 23 - Loss: 1.078 | Accuracy: 38.33%
```

```
julia> test_X = hcat(first.(test_data)...)
4×30 Matrix{Float64}:
-1.1392  0.0684325 -0.897674  0.551486  0.430722  -0.551486 -1.74302 -0.173094  0.551486 -0.776911
-1.27868 -0.131539 -1.27868  0.556746 -1.96696  0.786174 -0.131539 -0.590395 -0.590395  0.786174
 0.420326  0.760211 -0.429389  0.533621  0.420326  1.04345 -1.3924  0.193735  0.760211 -1.33575
 0.656838  0.788031 -0.130318  0.525645  0.394453  1.57519 -1.31105  0.132067  0.394453 -1.31105
```


```
julia> test_Y = hcat(last.(test_data)...)
3×30 OneHotMatrix{::Vector{UInt32}} with eltype Bool:
. . . 1 . 1 . . . . . 1 . . . 1 . . 1
. . 1 1 1 . 1 . . . . 1 . 1 1 1 1 1 . . 1 . 1 . . 1 .
1 1 . . . . . 1 1 1 1 . 1 . . . . 1 1 . . 1 1 . . 1 .
```

```
julia> println("Test Accuracy: ", round(accuracy(test_X, test_Y) * 100, digits=2), "%")
Test Accuracy: 16.67%
```


Training Loss



Training Accuracy



How Flux.jl Simplifies Deep Learning

Flux.jl provides a clean and intuitive way to define models using **Chain** and **Dense** layers. It supports automatic differentiation via Zygote.jl, ensuring seamless backpropagation. Additionally, GPU acceleration is built-in—just move the model and data to the GPU without modifying code.

Key Features

Automatic Differentiation (AD)

Definition: Automatic differentiation is a computational technique for efficiently and accurately evaluating derivatives (gradients) of functions. It works by decomposing functions into elementary operations and applying the chain rule.

Key Points in Flux.jl:

- Utilizes Zygote.jl for AD, performing source-to-source transformation to compute gradients.
- Supports dynamic AD—works seamlessly with native Julia control flow (e.g., loops, conditionals).
- Eliminates the need to manually construct a computational graph.
- Essential for training models using gradient descent, backpropagation, and optimization algorithms.

Benefits:

- Reduces human error in derivative calculations.
- Enables easy experimentation with complex model architectures.
- Supports differentiation through any differentiable Julia code.

GPU Acceleration

Definition: GPU acceleration offloads compute-intensive operations (such as matrix multiplications and convolutions) to the Graphics Processing Unit (GPU), which is optimized for parallel computation.

Key Points in Flux.jl:

- Leverages CUDA.jl, an interface for NVIDIA's CUDA API.
- Allows running the entire training loop (model, data, loss, gradients) on the GPU.
- Integrates seamlessly with Flux models—no extensive code modifications required.

Benefits:

- Significantly speeds up deep neural network training, especially for large datasets or complex models.
- Efficiently utilizes memory and computation resources on modern GPUs.
- Fully compatible with Julia, enabling high-performance computing in a streamlined environment.